

```

# Imports.
import numpy as np
import numpy.random as npr

from SwingyMonkey import SwingyMonkey

# also: keras for neural network
from keras.models import Sequential
from keras.layers.core import Dense, Dropout, Activation
from keras.optimizers import RMSprop, SGD

class Learner(object):
    '''
    This agent jumps randomly.
    '''

    def __init__(self):
        self.last_state = None
        self.last_action = None
        self.last_reward = None

        self.isFirstState = True # to catch weirdness for first state
of game
        self.isSecondState = False # update gravity on *second* state

        #self.Q = {} # initialize Q table. use a dictionary for now
        self.gamma = 1 # temporal discount value? finite horizon so
maybe we don't need this?
        #self.eps0 = 1 # do random action 5% of the time, for
exploration?
        self.eps = 0.2 # start more random
        self.g = 0
        self.alpha = 0.5

        ## initialize neural network to random weights
        # this should be optimized still
        self.model = Sequential()
        self.model.add(Dense(output_dim = 100,
batch_input_shape=(1,5), init = 'lecun_uniform' ))
        self.model.add(Activation("relu"))
        # two outputs for 2 actions!
        self.model.add(Dense(1, init='lecun_uniform'))
        self.model.add(Activation("linear")) #linear output so we can
have range of real-valued outputs

        ## initialize model
        rms = RMSprop()
        self.model.compile(loss='mse', optimizer=rms)

```

```

#sgd = SGD(lr=0.001)
#self.model.compile(loss='mean_squared_error', optimizer=sgd)

def reset(self, ii):
    self.last_state = None
    self.last_action = None
    self.last_reward = None
    self.g = 0
    # decrement epsilon value?
    #self.eps = float(self.eps0)/float(ii+1)
    self.eps = max([0.01, self.eps-0.001]) # always do at least
10% random actions
    self.isFirstState = True
    self.isSecondState = False

def getFeats(self, state, action):
    # takes state dictionary + action and converts to features
    # to feed into NN
    v = state['monkey']['vel']
    rx = state['tree']['dist']
    ry = state['monkey']['bot']-state['tree']['bot']
    h = state['monkey']['bot']

    # coarse-grain position here, if necessary?

    #dx = 1
    #dy = 1
    #rx = np.round(float(rx)/float(dx))
    #ry = np.round(float(ry)/float(dy))
    #h = np.round(float(h)/float(dy))

    #instead: normalize to max values?
    rx = float(rx)/float(300) # max dist of 300
    ry = float(ry)/float(200) # max diff here is +/- 200?
    v = float(v) / float(10) # guessing it's about 10ish max/min?
    g = float(self.g)/float(4) # 4 values?

    # can also coarse-grain velocity?
    #dv = 1
    #v = np.round(float(v)/float(dv))

    #tmp = [g, v, rx, ry, h] # 5-dimensional feature vector
    # now: try excluding height
    # also for now: ignore g?
    tmp = [g, v, rx, ry, action]
    # convert to 1x6 numpy array that NN expects
    feat = np.ndarray([1,5]) # 5 or however many features there
are
    feat[:] = tmp

```

```

        return feat

    def QNet(self, feat):
        # uses a neural network to approximate Q using a state-action
        pair (s,a)
        return self.model.predict(feat)

    def action_callback(self, state):
        """
        Implement this function to learn things and take actions.
        Return 0 if you don't want to jump and 1 if you do.
        """
        #print state
        # first, just sit tight if it's the first state
        if self.isFirstState:
            self.last_state = state
            self.last_action = 0
            #self.g = -state['monkey']['vel']
            #self.g = 1 # to train more quickly?
            # don't jump on the first state
            self.isFirstState = False
            self.isSecondState = True
            #self.model.train_on_batch(self.getFeats(self.last_state),
            0*np.random.rand(1,2))
            return 0

        # if second state, then update gravity
        if self.isSecondState:
            self.g = -state['monkey']['vel']
            self.isSecondState = False

        # You might do some learning here based on the current state
        and the last state.

        # You'll need to select an action and return it.
        # Return 0 to swing and 1 to jump.

        ## find Q values for old state
        Q0 = self.QNet(self.getFeats(self.last_state,
        self.last_action))
        # and for new state!
        Qstay = self.QNet(self.getFeats(state,0))
        Qjump = self.QNet(self.getFeats(state,1))
        # take max
        if Qjump > Qstay: # if Q value higher for jumping
            new_action = 1
            Qmax = Qjump
        else: # otherwise, don't jump

```

```

        new_action = 0
        Qmax = Qstay
        #print Qstay, Qjump, new_action

        # update target vector?
        Qtarg = self.last_reward + Qmax

        # gradient descent to update weights
        #self.model.fit(self.getFeats(self.last_state), Qtarg,
batch_size = 1, nb_epoch = 1, verbose = 0 )
        self.model.train_on_batch(self.getFeats(self.last_state,
self.last_action), Qtarg)

        # epsilon greedy: with probability epsilon, overwrite new
action w random one
        if npr.rand() < self.eps: # then choose randomly
            new_action = npr.rand() < 0.5

        # update last action and state
        self.last_action = new_action
        self.last_state = state

        # and return action
        return self.last_action

def reward_callback(self, reward):
    '''This gets called so you can see what reward you get.'''
    self.last_reward = reward

def run_games(learner, hist, iters = 100, t_len = 100):
    '''
    Driver function to simulate learning by having the agent play a
    sequence of games.
    '''

    for ii in range(iters):
        # Make a new monkey object.
        swing = SwingyMonkey(sound=False,                                # Don't
play sounds.
                                text="Epoch %d" % (ii),                # Display
the epoch on screen.
                                tick_length = t_len,                    # Make game
ticks super fast.
                                action_callback=learner.action_callback,
                                reward_callback=learner.reward_callback)

        # Loop until you hit something.
        while swing.game_loop():
            pass

```

```
# Save score history.
hist.append(swing.score)
print swing.score, learner.eps

# Reset the state of the learner.
learner.reset(ii)
#print learner.Q

return

if __name__ == '__main__':

    # Select agent.
    agent = Learner()

    # Empty list to save history.
    hist = []

    # Run games.
    run_games(agent, hist, 1000, 10)
    #print agent.Q

    # Save history.
    np.save('hist',np.array(hist))
```